

DATA 612 Assignment #4

Brian K. Liles

July 2, 2019

Goal

The goal of this assignment is give you practice working with accuracy and other recommender system metrics.

Deliverables

As in your previous assignments, compare the accuracy of at least two recommender system algorithms against your offline data.

Implement support for at least one business or user experience goal such as increased serendipity, novelty, or diversity.

Compare and report on any change in accuracy before and after you've made the change in #2.

As part of your textual conclusion, discuss one or more additional experiments that could be performed and/or metrics that could be evaluated only if online evaluation was possible. Also, briefly propose how you would design a reasonable online evaluation environment.

Libraries

Import Data

From the **recommenderlab** package, and with the guidance of the text **Building a Recommendation System with R** the data for use will be the **MovieLense**

```
data(MovieLense)
MovieLense
```

```
## 943 x 1664 rating matrix of class 'realRatingMatrix' with 99392 ratings.
```

```
# create a new variable entitled MSWebDF in order to take a look at the data
mlDF <- as(MovieLense, 'data.frame')
head(mlDF)
```

```
##      user                item rating
## 1      1      Toy Story (1995)      5
## 453    1      GoldenEye (1995)      3
## 584    1      Four Rooms (1995)     4
## 674    1      Get Shorty (1995)      3
## 883    1      Copycat (1995)        3
## 969    1 Shanghai Triad (Yao a yao yao dao waipo qiao) (1995)  5
```

```
# use the glimpse function to look at the newly created data frame
glimpse(mlDF)
```

```
## Observations: 99,392
## Variables: 3
## $ user    <fctr> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
## $ item    <fctr> Toy Story (1995), GoldenEye (1995), Four Rooms (1995), ...
## $ rating  <dbl> 5, 3, 4, 3, 3, 5, 4, 1, 5, 3, 2, 5, 5, 5, 5, 5, 3, 4, 5...
```

```
# summarize the rating variable
summary(mlDF$rating)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      1.00     3.00     4.00     3.53     4.00     5.00
```

Based off the findings from the **glimpse** and **summary** functions we see that the **MovieLense** dataset has 99,392 observations and 3 variables. In addition, the **ratings** column has a mean of 3.53, median of 4.00; data is slightly skewed to the left and scored on a scale from 1 to 5.

Similarity Matrix

```
similarity_users <- similarity(MovieLense[1:4,],
                              method = "cosine",
                              which = "users")
cat("Similarity Users Matrix Output","\n")
```


```
## Similarity Users Matrix Output
```

```
(similarityMatrix <- as.matrix(similarity_users))
```

```
##           1           2           3           4
## 1 0.0000000 0.9605820 0.8339504 0.9192637
## 2 0.9605820 0.0000000 0.9268716 0.9370341
## 3 0.8339504 0.9268716 0.0000000 0.9130323
## 4 0.9192637 0.9370341 0.9130323 0.0000000
```

User similarity can also be viewed using the **image** function

```
image(as.matrix(similarity_users), main = "MovieLense: Similarity of Users")
```



Recommender Models

The **recommenderlab** package has several algorithms which can create recommender models

```
recommender_models <- recommenderRegistry$get_entries(
  dataType = "realRatingMatrix")
names(recommender_models)
```

```
## [1] "ALS_realRatingMatrix"      "ALS_implicit_realRatingMatrix"
## [3] "IBCF_realRatingMatrix"    "POPULAR_realRatingMatrix"
## [5] "RANDOM_realRatingMatrix"   "RERECOMMEND_realRatingMatrix"
## [7] "SVD_realRatingMatrix"     "SVDF_realRatingMatrix"
## [9] "UBCF_realRatingMatrix"
```

According to the textbook, a rating equal to 0 represented a missing value. In addition, the **summary** function stated the column held values from 1 to 5; based off the text we will visualize the ratings

```
ratings <- as.vector(MovieLense@data)
cat("Table of MOvie Lense Ratings")
```

```
## Table of MOvie Lense Ratings
```

```
(table_ratings <- table(ratings))
```

```
## ratings
##      0      1      2      3      4      5
## 1469760  6059 11307 27002 33947 21077
```

When creating a recommender system, a potential customer would feel more comfortable with information from reliable sources. In the text, users who have rated at least 50 movies and watched 100 were used.

```
ratings_movies <- MovieLense[rowCounts(MovieLense) > 50,colCounts(MovieLense) > 100]
ratings_movies
```

```
## 560 x 332 rating matrix of class 'realRatingMatrix' with 55298 ratings.
```

Create Training Sets & Models for Comparison

Based off the **ratings_movies** we will build models based off k-fold to validate models. In the text a rating_threshold was set at 3, which is slightly under the mean which was 3.53.

```
test <- evaluationScheme(ratings_movies, method = "split", train = 0.8, k = 4, given = 15, goodRating = 3)

# method: this is the way to split the data
# train: this is the percentage of data in the training set
# given: number of items to keep
# goodRating: rating threshold
# k: number of times to run the evaluation
```

The next step is to evaluate the models using the IBCF, UBCF, ALS, POPULAR methods. In addition, the measures of accuracy will also be performed for each method: RMSE, MSE, MAE

IBCF

```
# IBCF Models
ibcf_recMod <- Recommender(getData(test, "train"), "IBCF")
ibcf_pred <- predict(ibcf_recMod, getData(test, "known"), type = "ratings")
cat("IBCF Method: RMSE, MSE, MAE", "\n", "\n")
```

```
## IBCF Method: RMSE, MSE, MAE
##
```

```
(ibcf <- calcPredictionAccuracy(ibcf_pred, getData(test, "unknown")))
```

```
##      RMSE      MSE      MAE
## 1.410003 1.988109 1.079129
```

```
ibcf_results <- evaluate(test, method = "IBCF", n = seq(10,100,10))
```

```
## IBCF run fold/sample [model time/prediction time]
##  1  [1.02sec/0.75sec]
##  2  [0.91sec/0.24sec]
##  3  [0.94sec/0.14sec]
##  4  [0.93sec/0.15sec]
```

```
class(ibcf_results)
```


```
## [1] "evaluationResults"
## attr(,"package")
## [1] "recommenderlab"
```

```
head(getConfusionMatrix(ibcf_results)[[1]])
```

```
##          TP          FP          FN          TN precision      recall          TPR
## 10  2.410714  7.589286  64.74107  242.2589  0.2410714  0.03346173  0.03346173
## 20  4.714286  15.285714  62.43750  234.5625  0.2357143  0.07155819  0.07155819
## 30  6.794643  23.205357  60.35714  226.6429  0.2264881  0.10386921  0.10386921
## 40  8.767857  31.232143  58.38393  218.6161  0.2191964  0.13313444  0.13313444
## 50 10.839286  39.160714  56.31250  210.6875  0.2167857  0.16319675  0.16319675
## 60 12.848214  47.151786  54.30357  202.6964  0.2141369  0.19429792  0.19429792
##          FPR
## 10  0.03010303
## 20  0.06095402
## 30  0.09288718
## 40  0.12529073
## 50  0.15694280
## 60  0.18916946
```


```
plot(ibcf_results, annotate = TRUE, main = "ROC Curve: IBCF Method")
```

ROC Curve: IBCF Method



```
plot(ibcf_results, "prec/rec", annotate = TRUE, main = "Precision-Recall: IBCF Method")
```

Precision-Recall: IBCF Method



ALS

```
# ALS Models
als_recMod <- Recommender(getData(test, "train"), "ALS")
als_pred <- predict(als_recMod, getData(test, "known"), type = "ratings")
cat("ALS Method: RMSE, MSE, MAE","\n","\n")
```

```
## ALS Method: RMSE, MSE, MAE
##
```

```
(als <- calcPredictionAccuracy(als_pred, getData(test, "unknown")))
```

```
##          RMSE          MSE          MAE
## 0.9533999 0.9089713 0.7562991
```

```
als_results <- evaluate(test, method = "ALS", n = seq(10,100,10))
```

```
## ALS run fold/sample [model time/prediction time]
##   1  [0.2sec/36.33sec]
##   2  [0sec/36sec]
##   3  [0sec/36.39sec]
##   4  [0sec/35.49sec]
```

```
class(als_results)
```


```
## [1] "evaluationResults"
## attr(,"package")
## [1] "recommenderlab"
```

```
head(getConfusionMatrix(als_results)[[1]])
```

```
##          TP          FP          FN          TN precision    recall      TPR
## 10  3.607143  6.392857 63.54464 243.4554 0.3607143 0.05652401 0.05652401
## 20  7.017857 12.982143 60.13393 236.8661 0.3508929 0.10737235 0.10737235
## 30 10.008929 19.991071 57.14286 229.8571 0.3336310 0.15066021 0.15066021
## 40 12.928571 27.071429 54.22321 222.7768 0.3232143 0.19177430 0.19177430
## 50 15.705357 34.294643 51.44643 215.5536 0.3141071 0.23220013 0.23220013
## 60 18.508929 41.491071 48.64286 208.3571 0.3084821 0.27514931 0.27514931
##          FPR
## 10 0.02508648
## 20 0.05078683
## 30 0.07833238
## 40 0.10619151
## 50 0.13461937
## 60 0.16311340
```

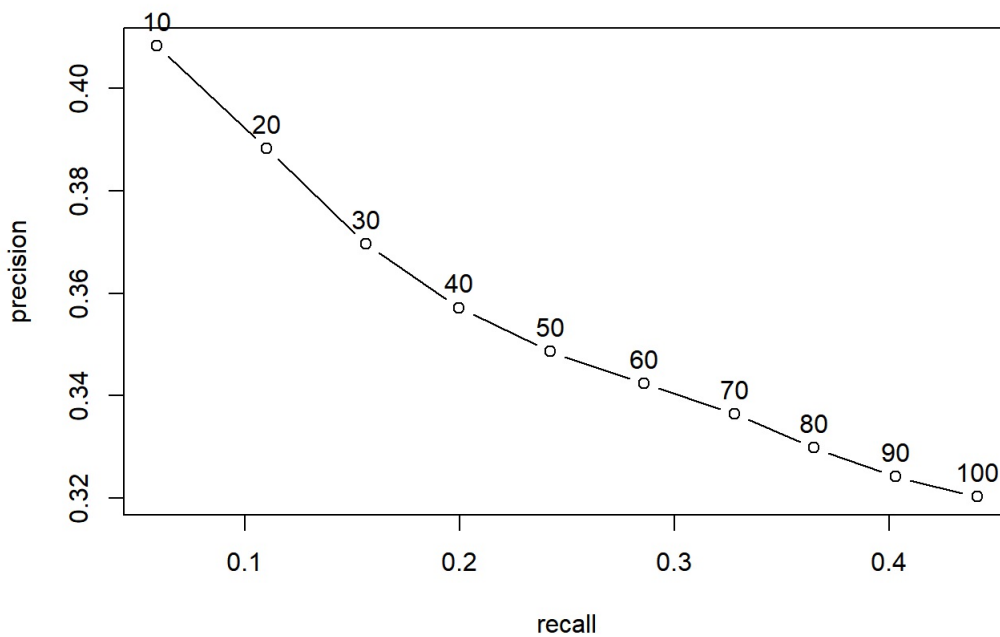
```
plot(als_results, annotate = TRUE, main = "ROC Curve: ALS Method")
```

ROC Curve: ALS Method



```
plot(als_results, "prec/rec", annotate = TRUE, main = "Precision-Recall: ALS Method")
```

Precision-Recall: ALS Method



UBCF

```
# UBCF Models
ubcf_recMod <- Recommender(getData(test, "train"), "UBCF")
ubcf_pred <- predict(ubcf_recMod, getData(test, "known"), type = "ratings")
cat("UBCF Method: RMSE, MSE, MAE", "\n", "\n")
```

```
## UBCF Method: RMSE, MSE, MAE
##
```

```
(ubcf <- calcPredictionAccuracy(ubcf_pred, getData(test, "unknown")))
```

```
##          RMSE          MSE          MAE
## 1.0279037 1.0565860 0.8206121
```

```
ubcf_results <- evaluate(test, method = "UBCF", n = seq(10,100,10))
```

```
## UBCF run fold/sample [model time/prediction time]
## 1 [0.1sec/0.58sec]
## 2 [0.02sec/0.56sec]
## 3 [0.02sec/0.6sec]
## 4 [0sec/0.59sec]
```

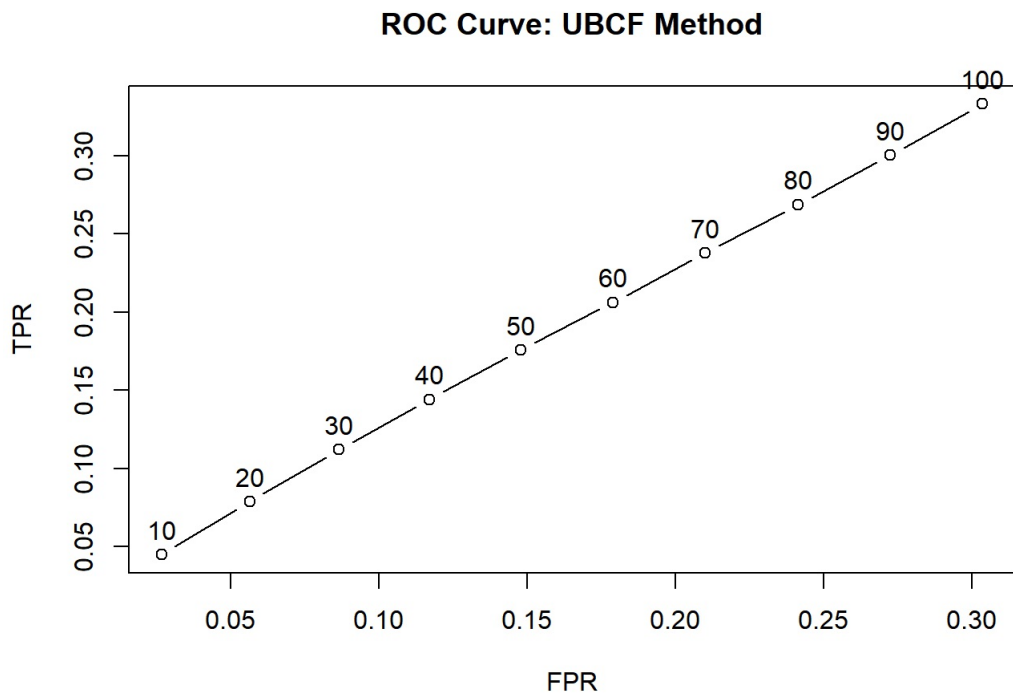
```
class(ubcf_results)
```

```
## [1] "evaluationResults"
## attr(,"package")
## [1] "recommenderlab"
```

```
head(getConfusionMatrix(ubcf_results)[[1]])
```


```
##           TP           FP           FN           TN precision      recall      TPR
## 10  3.205357  6.794643  63.94643  243.0536  0.3205357  0.04577616  0.04577616
## 20  5.464286  14.535714  61.68750  235.3125  0.2732143  0.07599436  0.07599436
## 30  7.812500  22.187500  59.33929  227.6607  0.2604167  0.11013581  0.11013581
## 40  9.982143  30.017857  57.16964  219.8304  0.2495536  0.14021526  0.14021526
## 50 12.241071  37.758929  54.91071  212.0893  0.2448214  0.17301485  0.17301485
## 60 14.312500  45.687500  52.83929  204.1607  0.2385417  0.20265402  0.20265402
##           FPR
## 10 0.02656125
## 20 0.05721992
## 30 0.08762384
## 40 0.11880667
## 50 0.14957927
## 60 0.18127923
```

```
plot(ubcf_results, annotate = TRUE, main = "ROC Curve: UBCF Method")
```



```
plot(ubcf_results, "prec/rec", annotate = TRUE, main = "Precision-Recall: UBCF Method")
```

Precision-Recall: UBCF Method



POPULAR

```
# POPULAR Models
popular_recMod <- Recommender(getData(test, "train"), "POPULAR")
popular_pred <- predict(popular_recMod, getData(test, "known"), type = "ratings")
cat("POPULAR Method: RMSE, MSE, MAE","\n","\n")
```

```
## POPULAR Method: RMSE, MSE, MAE
##
```

```
(popular <- calcPredictionAccuracy(popular_pred, getData(test, "unknown")))
```

```
##      RMSE      MSE      MAE
## 0.9720629 0.9449062 0.7634149
```

```
popular_results <- evaluate(test, method = "POPULAR", n = seq(10,100,10))
```

```
## POPULAR run fold/sample [model time/prediction time]
## 1 [0.09sec/0.58sec]
## 2 [0.02sec/0.61sec]
## 3 [0.01sec/0.57sec]
## 4 [0.03sec/0.56sec]
```

```
class(popular_results)
```

```
## [1] "evaluationResults"
## attr(,"package")
## [1] "recommenderlab"
```


```
head(getConfusionMatrix(popular_results)[[1]])
```


##	TP	FP	FN	TN	precision	recall	TPR
## 10	4.571429	5.428571	62.58036	244.4196	0.4571429	0.06788259	0.06788259
## 20	8.178571	11.821429	58.97321	238.0268	0.4089286	0.11811623	0.11811623
## 30	11.044643	18.955357	56.10714	230.8929	0.3681548	0.15977698	0.15977698
## 40	13.428571	26.571429	53.72321	223.2768	0.3357143	0.19384917	0.19384917
## 50	15.982143	34.017857	51.16964	215.8304	0.3196429	0.22985776	0.22985776
## 60	18.607143	41.392857	48.54464	208.4554	0.3101190	0.26653364	0.26653364

##	FPR
## 10	0.02059542
## 20	0.04526826
## 30	0.07332163
## 40	0.10339700
## 50	0.13277207
## 60	0.16174385


```
plot(popular_results, annotate = TRUE, main = "ROC Curve: POPULAR Method")
```

ROC Curve: POPULAR Method



```
plot(popular_results, "prec/rec", annotate = TRUE, main = "Precision-Recall: POPULAR Method")
```

Precision-Recall: POPULAR Method



```
cat("IBCF Model Time/Prediction Time Average","\n","\n")
```

```
## IBCF Model Time/Prediction Time Average
##
```

```
(ibcf_mean <- mean(c(0.15,0.24,0.14,0.14)))
```

```
## [1] 0.1675
```

```
cat("\n","ALS Model Time/Prediction Time Average","\n","\n")
```

```
##
## ALS Model Time/Prediction Time Average
##
```

```
(als_mean <- mean(c(37.38,36.89,37.8,37.5)))
```

```
## [1] 37.3925
```

```
cat("\n","UBCF Model Time/Prediction Time Average","\n","\n")
```

```
##
## UBCF Model Time/Prediction Time Average
##
```

```
(ubcf_mean <- mean(c(0.64,0.61,0.62,0.57)))
```

```
## [1] 0.61
```

```
cat("\n","POPULAR Model Time/Prediction Time Average","\n","\n")
```

```
##
## POPULAR Model Time/Prediction Time Average
##
```

```
(popular_mean <- mean(c(0.71,0.66,0.58,0.58)))
```

```
## [1] 0.6325
```

It is clear that the **IBCF** has the fastest model time/prediction time, while **ALS** is the slowest out of the 4 methods.

```
rbind(ibcf,als,ubcf,popular)
```

```
##           RMSE      MSE      MAE
## ibcf      1.410032 1.9881090 1.0791286
## als       0.9533999 0.9089713 0.7562991
## ubcf      1.0279037 1.0565860 0.8206121
## popular  0.9720629 0.9449062 0.7634149
```

Next, we will compare the **IBCF,UBCF** techniques

```
compare <- list(
  IBCF_cos = list(name = "IBCF", param = list(method = "cosine")),
  IBCF_cor = list(name = "IBCF", param = list(method = "pearson")),
  UBCF_cos = list(name = "UBCF", param = list(method = "cosine")),
  UBCF_cor = list(name = "UBCF", param = list(method = "pearson"))
)
```

```
compare_results <- evaluate(x = test, method = compare, n = c(1, 5, seq(10,100,10)))
```


```
## IBCF run fold/sample [model time/prediction time]
## 1 [0.99sec/0.15sec]
## 2 [0.98sec/0.16sec]
## 3 [0.98sec/0.14sec]
## 4 [1.13sec/0.15sec]
## IBCF run fold/sample [model time/prediction time]
## 1 [1.14sec/0.16sec]
## 2 [1.17sec/0.14sec]
## 3 [1.18sec/0.16sec]
## 4 [1.1sec/0.16sec]
## UBCF run fold/sample [model time/prediction time]
## 1 [0.02sec/0.59sec]
## 2 [0sec/0.61sec]
## 3 [0sec/0.59sec]
## 4 [0.01sec/0.6sec]
## UBCF run fold/sample [model time/prediction time]
## 1 [0.03sec/0.72sec]
## 2 [0.01sec/0.73sec]
## 3 [0.02sec/0.72sec]
## 4 [0.02sec/0.72sec]
```

```
class(compare_results)
```

```
## [1] "evaluationResultList"
## attr(,"package")
## [1] "recommenderlab"
```

```
plot(compare_results, annotate = 1, legend = "topleft")
title("ROC CURVE")
```

ROC CURVE



```
plot(compare_results, "prec/rec", annotate = 1, legend = "topleft")
title("PRECISION-RECALL")
```

PRECISION-RECALL

